

Partitioned Communication CUDA Bindings

Contact: Jim Dinan <jdinan@nvidia.com>

Approach

Define Kernel Triggered bindings for MPI partitioned operations

Approach is to define as part of MPI CUDA language bindings

- `__device__` APIs
- Callable by kernels (`__global__` functions)
 - `K<<<grid_dim, block_dim, smem, stream>>>(arg0, arg1, ...);`

Try to make them generic enough that similar bindings can be used with other accelerator languages / language extensions

MPI 4.0 Partitioned APIs

```
int MPI_Psend_init(const void *buf, int partitions, MPI_Count count,  
                  MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Info info,  
                  MPI_Request *request)
```

```
int MPI_Precv_init(void *buf, int partitions, MPI_Count count,  
                  MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Info info,  
                  MPI_Request *request)
```

```
int MPI_[start,wait][_all](...)
```

Keep host only

```
int MPI_Pready(int partition, MPI_Request request)
```

Add device bindings

```
int MPI_Pready_range(int partition_low, int partition_high, MPI_Request request)
```

```
int MPI_Pready_list(int length, const int array_of_partitions[], MPI_Request request)
```

```
int MPI_Parrived(MPI_Request request, int partition, int *flag)
```

MPI Request Objects in Device Context

MPI handle must be valid in the device context

- Internal MPI state must be accessible by device (e.g. device, pinned host memory)

Option 1: Unified Memory

- Allocate MPI object in managed memory, accessible to host/device, handle value is **unified** pointer
 - Pages containing MPI requests migrate on access by host/device, or can set `cudaMemAdviseSetPreferredLocation` to avoid migration
 - Alternatively, MPI can update requests in start call and then pin to the device via `cudaMemAdviseSetReadMostly`

Option 1.5: Unified Virtual Addressing

- Allocate MPI object in `cudaHostAlloc`'d memory, accessible to host/device, handle value is **unified** pointer
 - Host/device pointer are same when device supports `canUseHostPointerForRegisteredMem` property
 - Lives in pinned host memory, which adds PCIe access latency from GPU

Option 2: Export request handle to the device, handle value is **device** pointer

- New export operation exports MPI request so it can be accessed by device
- Improves portability and can optimize the object for access from the device

Request Export

```
__host__ int MPI_Prequest_create(MPI_Request req,  
                                MPI_Info info,  
                                MPI_Prequest *preq /* host ptr */)
```

```
__host__ int MPI_Prequest_free(MPI_Prequest *preq)
```

MPI_PREQUEST_NULL

- MPI_Prequest objects are only valid for use in device functions
- The MPI_Prequest object is created on the currently selected device
 - MPI_Request must be in the non-started state
 - Can optionally use the info argument to specify the device
- MPI_Prequest for an MPI_Psend operation may contain:
 - libibverbs: Pointer to list of pre-prepared RDMA write WQEs, pointer to QP (SQ, DBR, DB)
 - libfabric: Reference to event counter with associated triggered send operation
- MPI_Prequest for an MPI_Precv operation may contain:
 - Pointer to flag in memory
- Starting req also starts any preqs that were created on req
 - E.g. posts triggered operations to the NIC

Discussion 8-18-2021:

Ryan: Could add a device side request query function to convert a request object

```
__device__  
MPI_Request_query(  
    MPI_Request req_in,  
    MPI_Request *req_out);
```

- Import request to device
- Host creates “exportable” request object so it can be imported, not able to predict whether device will call query
- Could require memory allocation from device

Whit: Could make sense to have both prequest_create on host and query on the device

Device Ready

```
__device__ int MPI_Pready(MPI_Prequest prequest, int partition)
```

```
__device__ int MPI_Pready_range(int part_low, int part_high,  
                                MPI_Prequest prequest)
```

```
__device__ int MPI_Pready_list(int length,  
                                const int array_of_partitions[],  
                                MPI_Prequest prequest)
```

- MPI_Pready operation may involve
 - ibverbs: Copy pre-prepared WQE to SQ and ring doorbell
 - libfabric: Increment event counter
- Allow functions to be inlined
- Could replace `__device__` with a generic thing like `MPI_DEVICE_QUALIFIER`

Device Arrived

```
__device__ int MPI_Parrived(MPI_Prequest preq,  
                             int partition, int *flag)
```

- MPI_Parrived operation may involve
 - libibverbs: Polling a CQ, e.g. if write with immediate is used
 - Checking flags for completion of the given operation
- May generate more efficient code to return flag
- Allow functions to be inlined
- Q: Should we return an MPI error code?
 - What can the device do with it? Hand it back to the CPU.

Questions

Do we want warp/block variants of ready/arrived functions?

- E.g. may want to perform a memcpy in MPI_Pready for P2P transfers

Martin: Does this require MPI_THREAD_MULTIPLE or some other (new?) threading model in the MPI library?

- There is concurrency in two forms:
 - Multiple threads running on the device calling ready/arrived APIs
 - Host and device performing MPI operations at the same time
- Ryan: MPI_THREAD_PARTITIONED was part of the original partitioned communication proposal - Only parried/pready called by multiple threads, otherwise it's MPI_THREAD_SERIALIZED