

# Partition Communication and SYCL Bindings

Maria Garzaran



# A detour to SYCL



# oneAPI for Cross-Architecture Performance

Optimized Applications

Optimized Middleware & Frameworks

## oneAPI Product

Direct Programming

SYCL/DPC++

API-Based Programming

Libraries

Analysis &  
Debug Tools

CPU

GPU

FPGA

AI

Get functional quickly, then analyze and tune

# SYCL

- SYCL is an industry-driven Khronos standard that adds data parallelism to C++ for heterogeneous systems (<https://www.khronos.org/sycl/>)
- There are SYCL compilers that support different vendor's hardware:
  1. Intel compiler generates code for Intel CPUs, GPUs, and FPGAs
  2. ComputeCpp from CodePlay generates code for any CPU, Intel GPUs and FPGAs, AMD GPUs, ARM Mali, Nvidia GPUs
  3. triSYCL from Xilinx generates code Xilinx FPGAs
  4. hipSYCL from University of Heidelberg generates code for AMD GPUs and Nvidia GPUs.
- Libraries such as Kokkos and Raja have SYCL backends

# Hello data-parallel SYCL code

```
1. #include <CL/sycl.hpp>
2. #include <iostream>
3. using namespace sycl;
4.
5. const std::string secret {
6.     "Ifmmp-!xpsme!\012J(n!tpssz-!Ebwf/!"
7.     "J(n!bgsbje!J!dbo(u!ep!uibu/!..!BM\01" };
8. const auto sz = secret.size();
9.
10. int main() {
11.     queue Q;
12.
13.     char *result = malloc_shared<char>(sz, Q);
14.     std::memcpy(result, secret.data(), sz);
15.
16.     Q.parallel_for(sz, [=](auto& i) {
17.         result[i] -= 1;
18.     }).wait();
19.
20.     std::cout << result << "\n";
21.     return 0;
22. }
```

- Work is submitted to a queue
- A queue is associated with a single device
- Work in the device is asynchronous

- Allocate memory on shared memory
- (both host and device can access)

- Submit work to the queue
- The work is expressed as a lambda function here

- Piece of work the kernel performs

# Hello data-parallel SYCL code

```
1. #include <CL/sycl.hpp>
2. #include <iostream>
3. using namespace sycl;
4.
5. const std::string secret {
6.     "lfmmp-!xpsme\"012J(n!tpssz-!Ebwf/!"
7.     "J(n!bgsbje!J!dbo(u!ep!uibu/!.!IBM\01" };
8. const auto sz = secret.size();
9.
10. int main() {
11.     queue Q;
12.
13.     char *result = malloc_shared<char>(sz, Q);
14.     std::memcpy(result, secret.data(), sz);
15.
16.     Q.parallel_for(sz, [=](auto& i) {
17.         result[i] -= 1;
18.     }).wait();
19.
20.     std::cout << result << "\n";
21.     return 0;
22. }
```

Host code

Device code

Host code

Where the code executes depends on where we setup the queue

```
queue Q { default_selector{} };
queue Q { host_selector{} };
queue Q { cpu_selector{} };
queue Q { gpu_selector{} };
queue Q { accelerator_selector{} };
queue Q { INTEL::fpga_selector{} };
```

# Partition Communication



## Example 4.2 MPI standard (OpenMP)

```
if (myrank == 0) {  
    MPI_Psend_init(message, partitions, count, xfer_type,  
        dest, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    #pragma omp parallel for shared(request)  
    num_threads(NUM_THREADS)  
    for (int i = 0; i < partitions; i++) {  
        /* compute and fill partition #i, then mark ready: */  
        MPI_Pready(i, request);  
    }  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

```
if (myrank == 1) {  
    MPI_Precv_init(message, partitions, count, xfer_type,  
        source, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```



## Partition Communication Ex. 4.2 MPI Standard (SYCL)

```
if (myrank == 0) {  
    MPI_Psend_init(message, partitions, count, xfer_type,  
        dest, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    mQueue.submit([&](sycl::handler &h) {  
        h.parallel_for(partitions, [=](id<1> i) {  
            MPI_Pready(i, request);  
        });  
    });  
    mQueue.wait();  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

```
if (myrank == 1) {  
    MPI_Precv_init(message, partitions, count, xfer_type,  
        source, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

# Partition Communication: MPI\_Pready executes from Device

```
if (myrank == 0) {  
    MPI_Psend_init(message, partitions, count, xfer_type,  
        dest, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
    MPI_Pbuf_prepare(&request);  
  
    mQueue.submit([&](sycl::handler &h) {  
        h.parallel_for(partitions, [=](id<1> i) {  
            MPI_Pready(i, request);  
        });  
    });  
    mQueue.wait();  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

```
if (myrank == 1) {  
    MPI_Precv_init(message, partitions, count, xfer_type  
        source, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
    MPI_Pbuf_prepare(&request);  
  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

# Partition Communication: MPI\_Prequest for device

```
if (myrank == 0) {  
    MPI_Psend_init(message, partitions, count, xfer_type,  
        dest, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
    MPI_Pbuf_prepare(&request);  
    // Create a request to pass to the kernel  
    MPI_Prequest_create(request, info, &device_request);  
  
    mQueue.submit([&](sycl::handler &h) {  
        h.parallel_for(partitions, [=](id<1> i) {  
            MPI_Pready(i, device_request);  
        });  
    });  
    mQueue.wait();  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

1/19/22

```
if (myrank == 1) {  
    MPI_Precv_init(message, partitions, count, xfer_type,  
        source, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
    MPI_Pbuf_prepare(&request);  
  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

# Partition Communication: MPI\_Prequest for device (notes)

```
if (myrank == 0) {
    MPI_Psend_init(message, partitions, count, xfer_type,
        dest, tag, info, MPI_COMM_WORLD, &request);
    MPI_Start(&request);
    MPI_Pbuf_prepare(&request);
    // Create a request to pass to the kernel
    MPI_Prequest_create(request, info, &device_request);

    mQueue.submit([&](sycl::handler &h) {
        h.parallel_for(partitions, [=](id<1> i) {
            MPI_Pready(i, device_request);
        });
    });
    mQueue.wait();
    while (!flag) {
        /* Do useful work */
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);
    }
    MPI_Request_free(&request);
}
```

1/19/22

```
if (myrank == 1) {
    MPI_Precv_init(message, partitions, count, xfer_type
        source, tag, info, MPI_COMM_WORLD, &request);
    MPI_Start(&request);
    MPI_Pbuf_prepare(&request);
```

- We are making the implicit assumption that the device request is allocated on the same device as the message
- We also assume that the message is only allocated in a single device
- Device information is not needed, because we assume the runtime is able to recover the device where message is allocated from the request created during the MPI\_Psend\_init() call
- An additional reason to have the device\_request is that in some devices (e.g. FPGAs), the kernel might not be able to access the host memory, so it would not be able to access the request in the host memory

```
    }
    MPI_Request_free(&request);
}
```

# MPI\_Prequest

```
int MPI_Prequest_create (MPI_Request req,  
                        MPI_Info info,  
                        MPI_Prequest *preq)  
  
int MPI_Prequest_free(MPI_Prequest *preq)
```

Executes in the host

- MPI\_Prequest should be created on the same device as the message in MPI\_Send\_Init().
- The MPI implementation is free to allocate the MPI\_Prequest on USM shared, device or host.
- Each MPI implementation can determine the information that is needed on the MPI\_Prequest
- MPI\_Prequest objects are only valid for use in device functions

# Device Ready

SYCL\_EXTERNAL int MPI\_Pready(MPI\_Prequest prequest, int partition)

SYCL\_EXTERNAL int MPI\_Pready\_range(int part\_low, int part\_high, MPI\_Prequest prequest)

SYCL\_EXTERNAL int MPI\_Pready\_list(int length, const int array\_of\_partitions[], MPI\_Prequest prequest)

Execute in the device

# Device Arrived

```
SYCL_EXTERNAL int MPI_Parrived(MPI_Prequest preq, int partition, int *flag)
```

Execute in the device

# Others

- For efficiency reasons we should define variants of MPI\_Ready()/MPI\_Parrived() for work items (warp) and/or work groups (thread block)



## Example 4.2 MPI standard (OpenMP offload)

```
if (myrank == 0) {  
    MPI_Psend_init(message, partitions, count, xfer_type,  
        dest, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    #pragma omp target teams distribute parallel for  
    for (int i = 0; i < partitions; i++) {  
        /* compute and fill partition #i, then mark ready: */  
        MPI_Pready(i, request);  
    }  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

```
if (myrank == 1) {  
    MPI_Precv_init(message, partitions, count, xfer_type,  
        source, tag, info, MPI_COMM_WORLD, &request);  
    MPI_Start(&request);  
  
    while (!flag) {  
        /* Do useful work */  
        MPI_Test(&request, &flag, MPI_STATUS_IGNORE);  
    }  
    MPI_Request_free(&request);  
}
```

# OpenMP declare target directive

**declare target** directive is used to indicate that the corresponding call inside a **target** region is to a function that can execute on the default target device

```
#pragma omp declare target
```

```
int MPI_Pready(MPI_Request request, int partition) {
```

```
...
```

```
}
```

```
#pragma omp end declare target
```

