

# Redefining the Future of Accelerator Computing with Level Zero

Presenter: Jaime Arteaga

Contributors: Michal Mrozek, Ravindra Babu Ganapathi, Ben Ashbaugh, Brandon Fliflet, Aravind Gopalakrishnan, and Maria Garzaran

Hybrid Working Group MPI Forum 2022

The Intel logo is located in the bottom left corner. It consists of the word "intel" in a white, lowercase, sans-serif font, followed by a registered trademark symbol (®). To the left of the text is a graphic element made of several blue squares of varying sizes, some of which are slightly offset from each other, creating a pixelated or mosaic-like effect.

intel®

# Outline

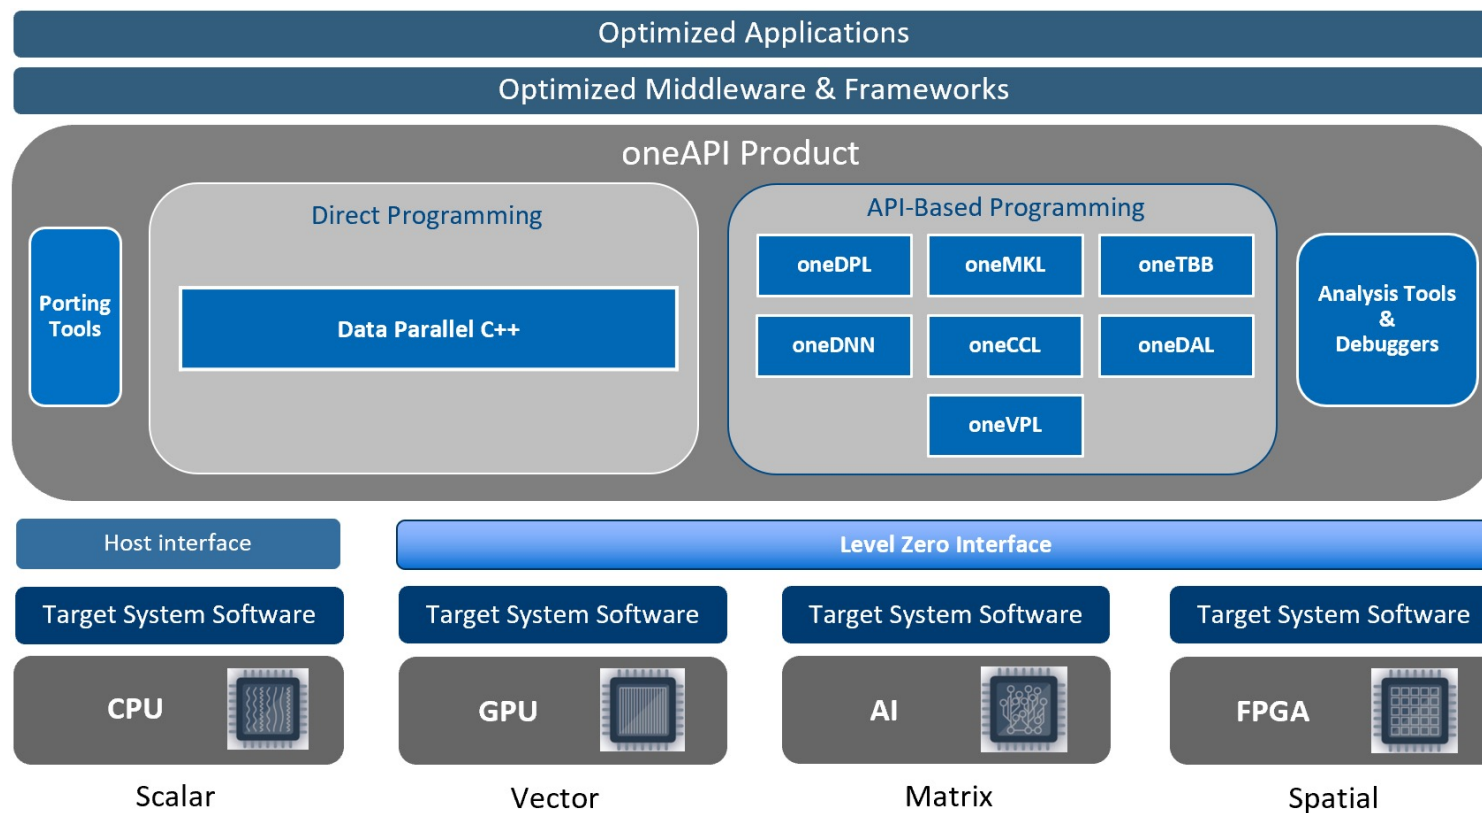
- *What* is Level Zero?
- *Level Zero Device Model*
- *Level Zero Memory Model*
- *Level Zero Scheduling Model*
- *IPC*
- *Summary*

# What is Level Zero?

- Low-level driver interface exposing underlying device capabilities to higher level languages
- Abstraction layer for HW accelerators that can be implemented by any vendor for any type of accelerator
- More control and lower-level access to rich device feature set
  - Leads to higher performance and functionality
  - Low latency direct-to-metal interface
  - Support many core threaded applications (e.g., batching)
- Support for broader language features such as
  - Function pointers
  - Unified memory
  - I/O capabilities

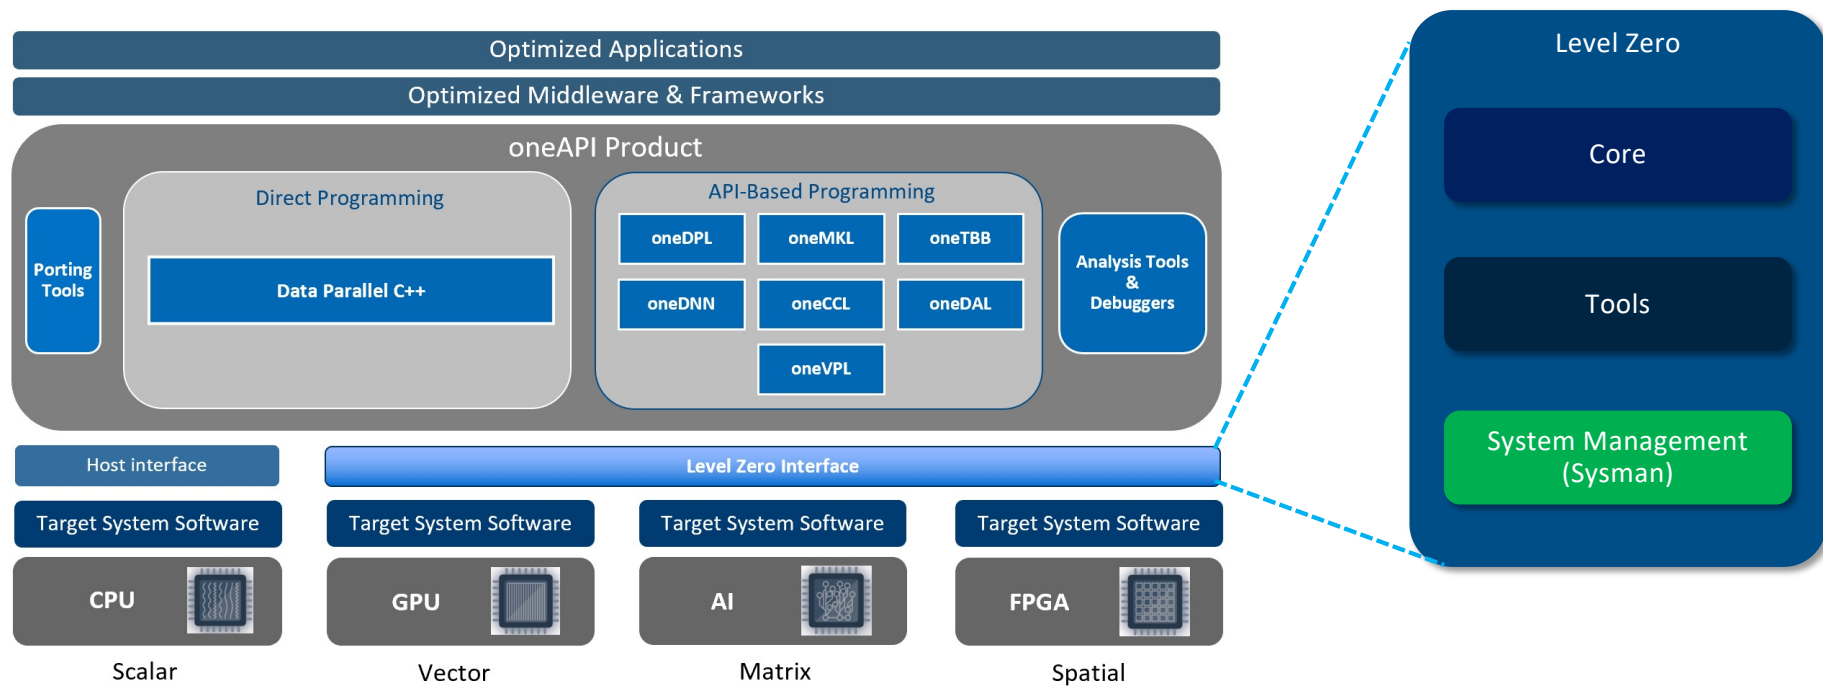
<https://spec.oneapi.io/level-zero/latest/index.html>

# oneAPI and Level Zero Software Stack

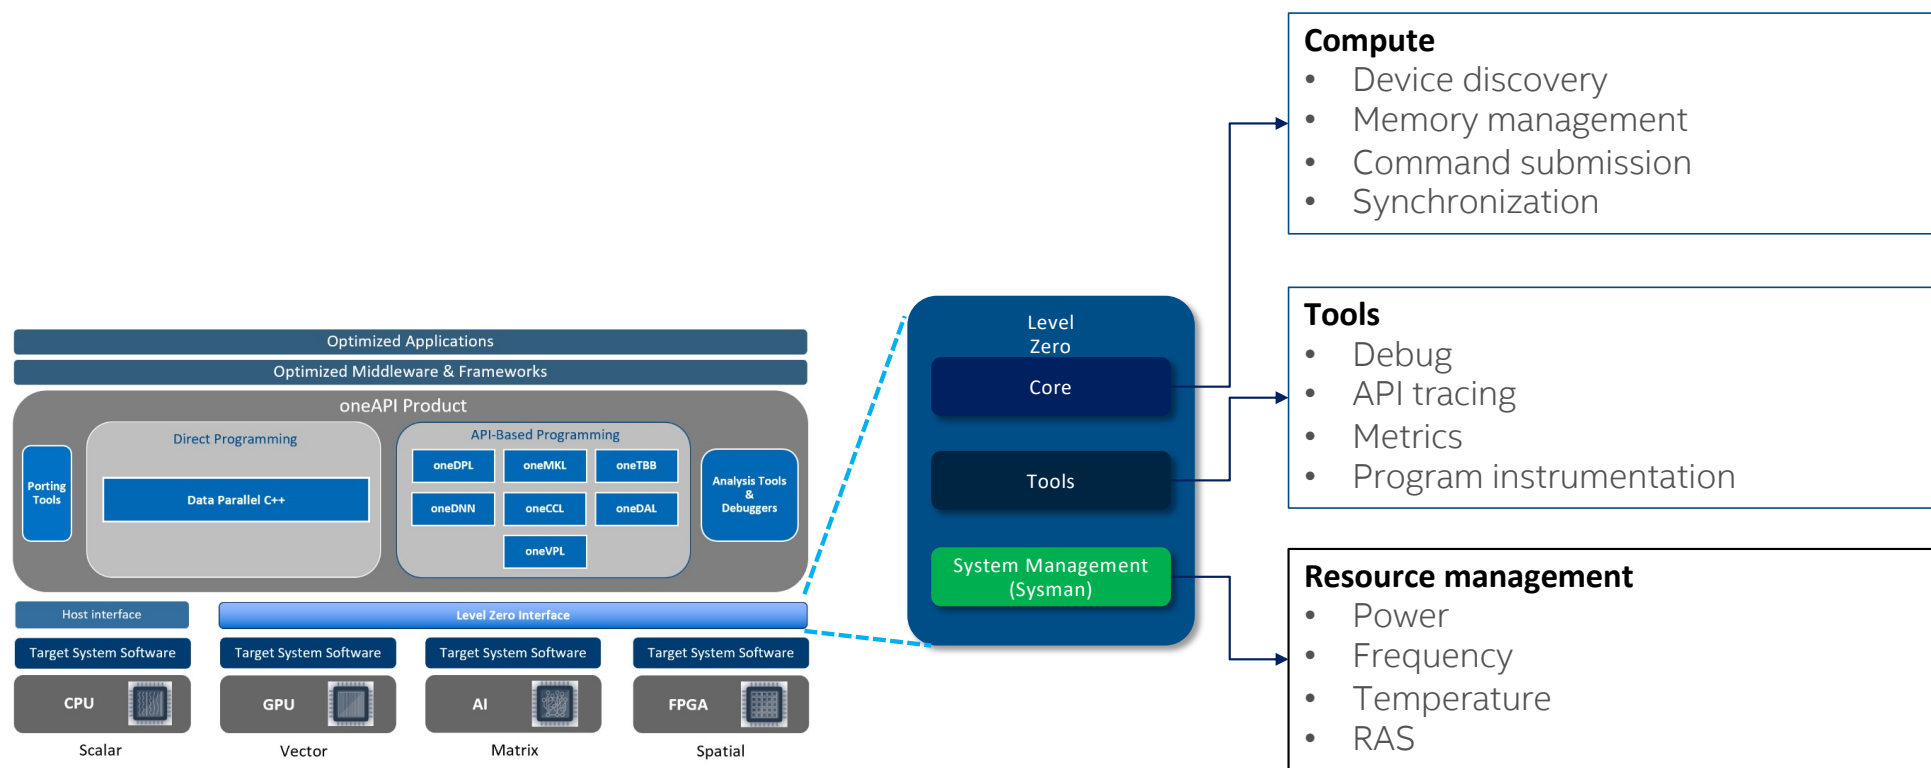


<https://spec.oneapi.io/level-zero/latest/core/INTRO.html#objective>

# oneAPI and Level Zero Software Stack - I



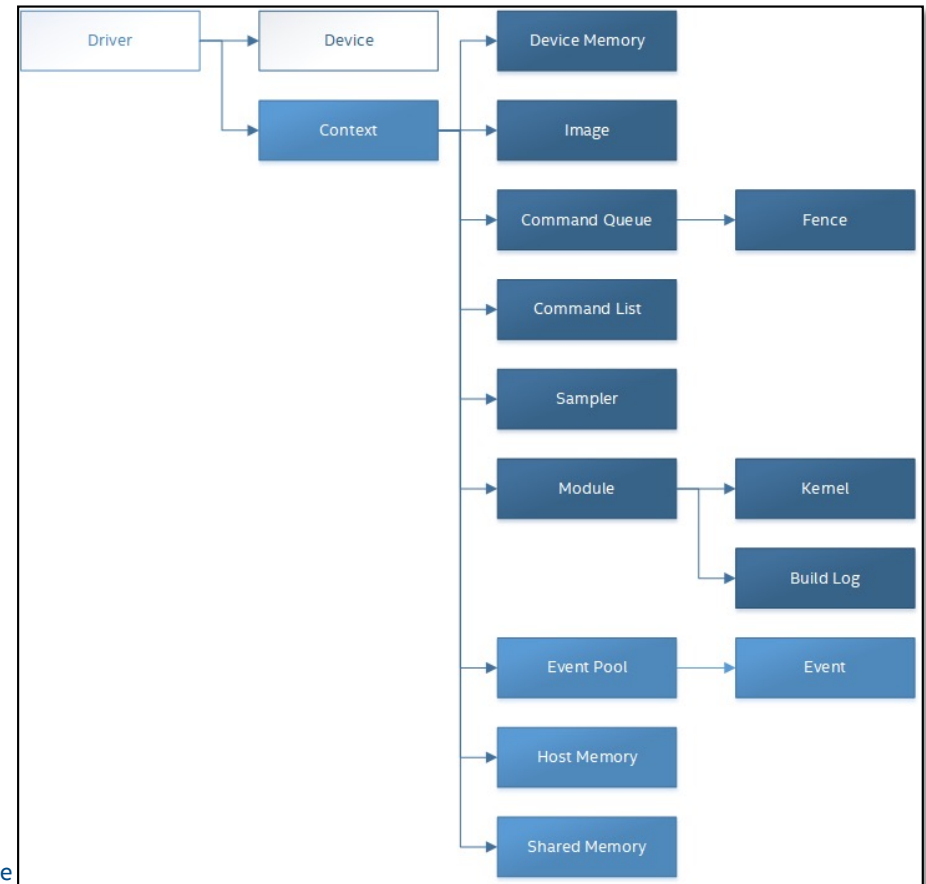
# oneAPI and Level Zero Software Stack - II



# Level Zero Drivers and Contexts

- Driver objects represent a collection of physical devices in system
- More than one driver may be available in the system.
  - Example: One driver may support two accelerators from one vendor and another driver support an accelerator from a different vendor.
- Contexts are primarily used during creation and management of resources that may be used by multiple devices.

<https://spec.oneapi.io/level-zero/latest/core/PROG.html#device>



# Level Zero Device Discovery

- Device object represents a physical device in the system
- Device discovery API enumerates devices in system
  - Use `zeDeviceGet` to
    - Query number of Level Zero devices supported by driver
    - Obtain devices objects which are read-only global constructs.
  - Device Properties queried using `zeDeviceGetProperties`
    - Device type, max memory allocation size, ...
  - Universal Unique Identifier (UUID)
    - 16-byte globally unique and immutable identifier
    - Uniquely identify particular device in a node within a datacenter
- Device handle is primarily used during creation and management of resources that are specific to a device

<https://spec.oneapi.io/level-zero/latest/core/PROG.html#drivers-and-devices>



# Sub-Devices

- Level Zero allows for fine-grain abstraction of HW by exposing *sub-devices*
- Implementations are free to define what a sub-device is, including:
  - Compute capabilities: Number of queues, scheduling policies available.
  - Memory: How much memory each sub-device uses.
  - Memory affinity: How memory affinitize to available queues.
- API available to query and obtain a sub-device from a parent device
- No distinction between sub-devices and parent devices in scheduling and allocation interfaces, as both use same handle (or data type).

# Level Zero Device Discovery

L0 driver initialization

Driver discovery

Context creation

Device discovery

```
zeInit(ZE_INIT_FLAG_GPU_ONLY);

uint32_t driverCount = 0;
zeDriverGet(&driverCount, nullptr);
if (driverCount == 0) {
    std::terminate();
}

zeDriverGet(&driverCount, &driverHandle);

ze_context_desc_t contextDesc = {};
zeContextCreate(driverHandle, &contextDesc, &context);

uint32_t deviceCount = 0;
zeDeviceGet(driverHandle, &deviceCount, nullptr);
if (deviceCount == 0) {
    std::terminate();
}

std::vector<ze_device_handle_t> devices(deviceCount);
zeDeviceGet(driverHandle, &deviceCount, devices.data());
```

# Level Zero Memory

- Visible to upper-level software stack as unified memory with single VA space.
- Two types of allocations
  - Memory – linear, unformatted allocations for direct access from host/device
  - Images – non-linear, formatted allocations for direct access from device
- Memory
  - Host – Owned by host and accessible by host and one or more devices
  - Device – Owned by device (local mem) and generally not accessible by host.
  - Shared – Share ownership between device and host (intended to migrate)
- Designed support for system allocator (e.g., malloc/new)

<https://spec.oneapi.io/level-zero/latest/core/PROG.html#memory-and-images>

# Level Zero Memory

Host Memory

```
void *hostBuffer = nullptr;  
ze_host_mem_alloc_desc_t hostDesc = {ZE_STRUCTURE_TYPE_HOST_MEM_ALLOC_DESC};  
zeMemAllocHost(context, &hostDesc, srcMemorySize, 1, &hostBuffer);
```

Device Memory

```
void *deviceBuffer = nullptr;  
ze_device_mem_alloc_desc_t deviceDesc = {ZE_STRUCTURE_TYPE_DEVICE_MEM_ALLOC_DESC};  
zeMemAllocDevice(context, &deviceDesc, allocSize, allocSize, device, &deviceBuffer);
```

Shared Memory

```
void *sharedBuffer = nullptr;  
ze_device_mem_alloc_desc_t deviceDesc = {ZE_STRUCTURE_TYPE_DEVICE_MEM_ALLOC_DESC};  
ze_host_mem_alloc_desc_t hostDesc = {ZE_STRUCTURE_TYPE_HOST_MEM_ALLOC_DESC};  
zeMemAllocShared(context, &deviceDesc, &hostDesc, allocSize, 1, device, &sharedBuffer);
```

# Level Zero Scheduling Model

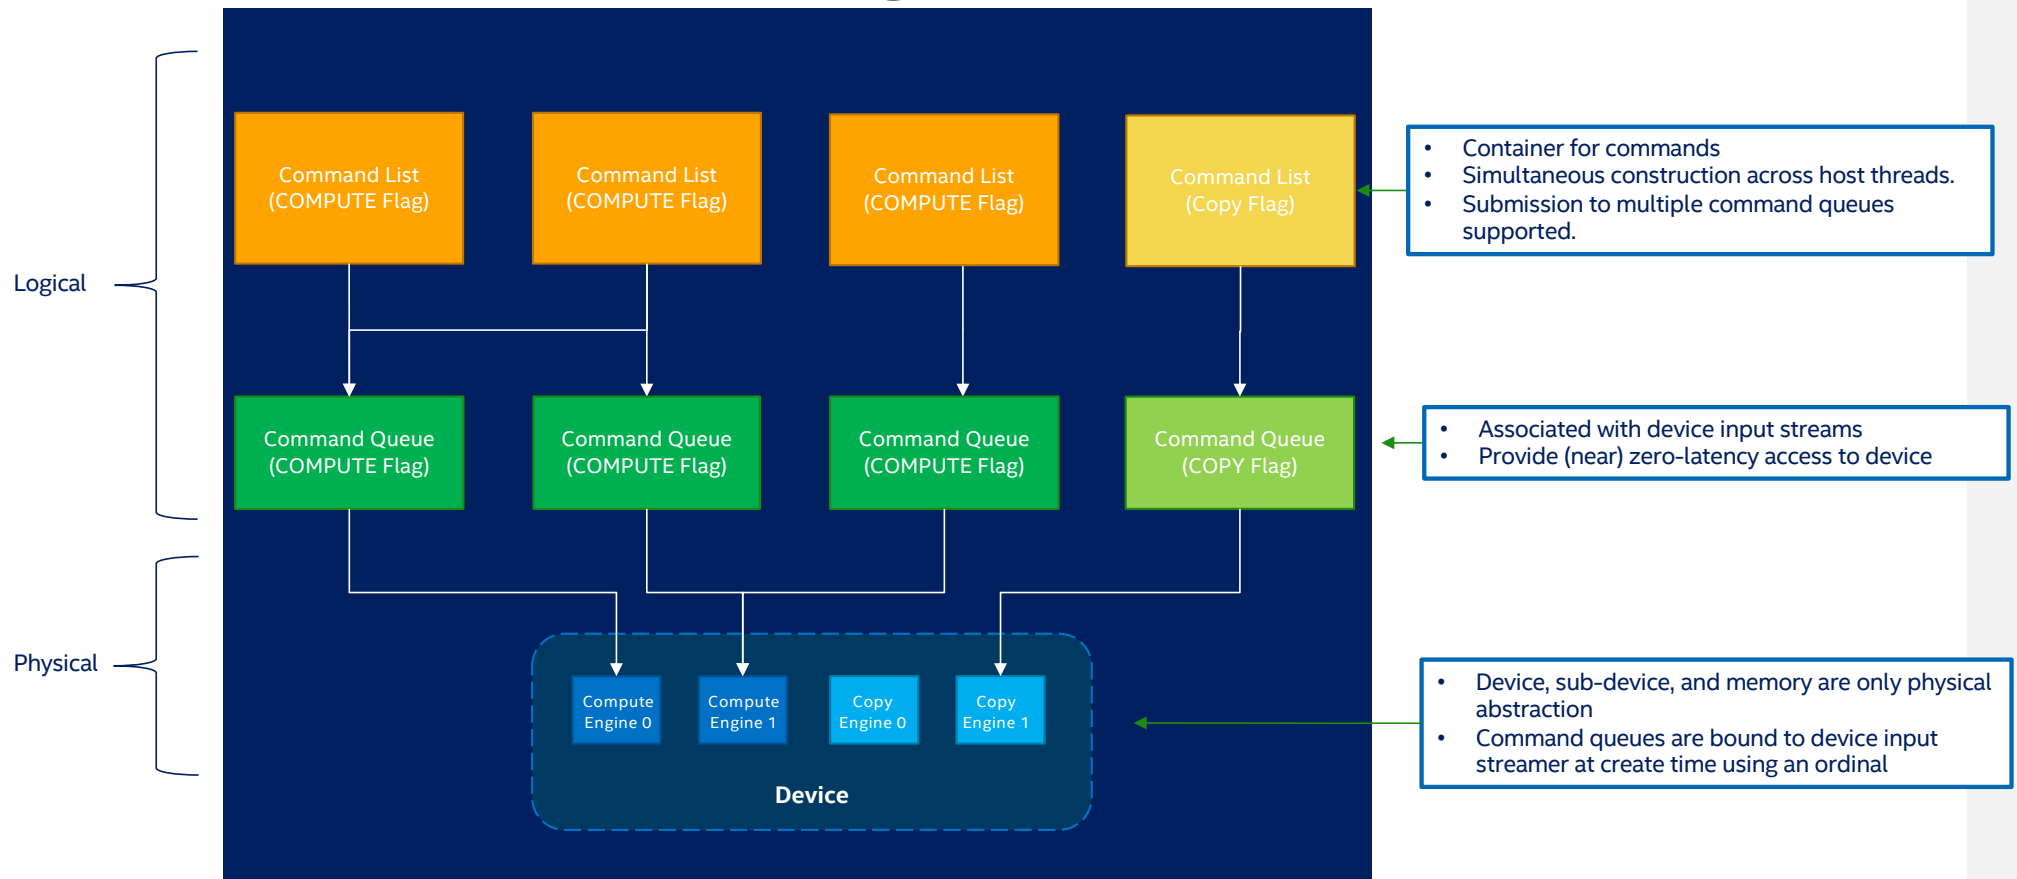
- Commands (user kernels, copies, synchronization actions) are appended to **command lists**.
  - A command list represents a sequence of commands to be executed in the accelerator
    - Can be recycled by resetting the list, without needing to create it again.
    - Can be reused, by submitting the same sequence the commands several times, without needing to re-append commands.
- Command lists are then submitted for execution to a **command queue**.
  - Logical object associated to a physical input stream in the device
  - Can be configured as synchronous or asynchronous
  - Types of command queues exposed as **Queue Groups**
    - **Compute** and **Copy** groups in a GPU for instance

<https://spec.oneapi.io/level-zero/latest/core/PROG.html#command-queues-and-command-lists>

# Level Zero Scheduling Model

- Commands and submissions can be synchronized using:
  - **Events:** Fine-grain synchronization between commands, host, and devices
    - Can be associated with any command list
    - Can be shared between processes and devices
    - Can be signaled/wait upon on either host or device, synchronously or asynchronously
    - Can provide timestamps for kernel execution
  - **Fences:** Coarse-grain synchronization of a command queue submission, signaled by device and waited upon on the host

# Level Zero Scheduling Model



<https://spec.oneapi.io/level-zero/latest/core/PROG.html#command-queues-and-command-lists>

# Level Zero Scheduling Model

Query queue groups

Look for a queue group with compute capabilities

Create command queue

Create command list

```
zeCall(zeDeviceGetCommandQueueGroupProperties(device, &numQueueGroups, nullptr));

std::vector<ze_command_queue_group_properties_t> queueProperties(numQueueGroups);
zeCall(zeDeviceGetCommandQueueGroupProperties(device, &numQueueGroups,
                                             queueProperties.data()));

uint32_t computeOrdinal = std::numeric_limits<uint32_t>::max();
for (uint32_t i = 0; i < numQueueGroups; i++) {
    if (queueProperties[i].flags & ZE_COMMAND_QUEUE_GROUP_PROPERTY_FLAG_COMPUTE) {
        computeOrdinal = i;
        break;
    }
}

ze_command_queue_handle_t queue;
ze_command_queue_desc_t cmdQueueDesc = {};
cmdQueueDesc.ordinal = computeOrdinal;
cmdQueueDesc.mode = ZE_COMMAND_QUEUE_MODE_ASYNCHRONOUS;
zeCommandQueueCreate(context, device, &cmdQueueDesc, &queue);

ze_command_list_handle_t list;
ze_command_list_desc_t listDesc = {};
listDesc.commandQueueGroupOrdinal = ordinal;
zeCommandListCreate(context, this->device, &listDesc, &list);
```



# Level Zero Scheduling Model

Module creation

Kernel creation

Kernel setup

Append kernel

Submit kernel for execution

Wait for completion

```
ze_module_handle_t module = nullptr;
ze_module_desc_t moduleDesc = {ZE_STRUCTURE_TYPE_MODULE_DESC};
ze_module_build_log_handle_t buildlog;
moduleDesc.format = ZE_MODULE_FORMAT_IL_SPIRV;
moduleDesc.pInputModule = reinterpret_cast<const uint8_t *>(spirvInput.get());
moduleDesc.inputSize = length;
moduleDesc.pBuildFlags = "";
zeModuleCreate(context, device, &moduleDesc, &module, nullptr);

ze_kernel_handle_t kernel = nullptr;
ze_kernel_desc_t kernelDesc = {ZE_STRUCTURE_TYPE_KERNEL_DESC};
kernelDesc.pKernelName = "CopyBufferToBufferBytes";
zeKernelCreate(module, &kernelDesc, &kernel);

uint32_t groupSizeX = 32u;
uint32_t groupSizeY = 1u;
uint32_t groupSizeZ = 1u;
zeKernelSuggestGroupSize(kernel, allocSize, 1U, 1U, &groupSizeX, &groupSizeY, &groupSizeZ);
zeKernelSetGroupSize(kernel, groupSizeX, groupSizeY, groupSizeZ);

zeKernelSetArgumentValue(kernel, 1, sizeof(dstBuffer), &dstBuffer);
zeKernelSetArgumentValue(kernel, 0, sizeof(srcBuffer), &srcBuffer);

ze_group_count_t dispatchTraits;
dispatchTraits.groupCountX = allocSize / groupSizeX;
dispatchTraits.groupCountY = 1u;
dispatchTraits.groupCountZ = 1u;

zeCommandListAppendLaunchKernel(cmdList, kernel, &dispatchTraits,
                                nullptr, 0, nullptr);

zeCommandListClose(cmdList);
zeCommandQueueExecuteCommandLists(cmdQueue, 1, &cmdList, nullptr);

zeCommandQueueSynchronize(cmdQueue, std::numeric_limits<uint64_t>::max());
```

# Inter-Process Communication (IPC)

- IPC calls allow sharing of memory objects across different device processes.
- There are two types of IPC APIs across processes:
  - Memory
  - Events

# Inter-Process Communication (Memory)

- The sender process:
  - Gets a handle for the allocation on the sending process through `zeMemGetIpcHandle`
  - Sends the handle to the receiving process
- The receiver process
  - Receive the handle from receiving process
  - Opens the handle on the virtual space of the receiving process through `zeMemOpenIpcHandle`

# Inter Process Communication (Memory)

Memory allocation

GetIpcHandle

Send Ipc Handle

Receive Ipc Handle

Open Ipc Handle

```
void* dptr = nullptr;
zeMemAllocDevice(hContext, &desc, size, alignment, hDevice, &dptr);

ze_ipc_mem_handle_t hIPC;
zeMemGetIpcHandle(hContext, dptr, &hIPC);

send_to_receiving_process(hIPC);
```

```
ze_ipc_mem_handle_t hIPC;
hIPC = receive_from_sending_process();

void* dptr = nullptr;
zeMemOpenIpcHandle(hContext, hDevice, hIPC, 0, &dptr);
```

sender  
rank

receiver  
rank

A method to send the Ipc handle is shown in [https://github.com/TApplencourt/GPU\\_IPC\\_Handle/blob/main/syclo\\_ipc\\_copy\\_dma\\_buf.cpp#L22](https://github.com/TApplencourt/GPU_IPC_Handle/blob/main/syclo_ipc_copy_dma_buf.cpp#L22)

Or you can use the pidfd\_getfd system call: <https://github.com/intel/compute-runtime/issues/448>

# Level Zero Websites

- Full open-source!
  - Specification:
    - <https://spec.oneapi.io/level-zero/latest/index.html>
  - Level Zero loader:
    - <https://github.com/oneapi-src/level-zero>
  - Level Zero Conformance Tests:
    - <https://github.com/oneapi-src/level-zero-tests>
  - Level Zero Implementation for Intel GPUs:
    - <https://github.com/intel/compute-runtime>

# Programming models already using Level Zero

- DPC++ Level Zero Plugin:
  - [https://github.com/intel/llvm/tree/sycl/sycl/plugins/level\\_zero](https://github.com/intel/llvm/tree/sycl/sycl/plugins/level_zero)
- oneAPI Julia
  - <https://github.com/JuliaGPU/oneAPI.jl>
- Intel oneAPI Level Zero Introduction
  - <https://www.intel.com/content/www/us/en/develop/documentation/oneapi-dpcpp-cpp-compiler-dev-guide-and-reference/top/optimization-and-programming-guide/intel-oneapi-level-zero-introduction.html>

# Summary

- Level Zero provides a rich set of interfaces to schedule work and manage memory on different accelerators.
- Objects defined in Level Zero, such as Command Queues and Command Lists, allow for a low-level control of the underlying hardware.
- With these and available optimization techniques, high-level programming languages and applications may execute workloads with close-to-metal latencies for higher performance.

